

Université Mohammed VI Polytechnique

College of Computing

TP4

Report

Module: Foundations of Parallel Computing

Instructor: Imad Kissami

Prepared by:

Malak Kably

Academic Year: 2025–2026

Contents

1 Exercise 1: Work Distribution with Parallel Sections	2
1.1 Implementation	2
1.1.1 Modified Code	2
1.2 Justification	2
2 Exercise 2: Master and Single Thread Directives	3
2.1 Implementation	3
2.1.1 Modified Code	3
2.2 Compilation and Execution	3
2.3 Results Comparison	3
3 Exercise 3: Load Balancing with Parallel Sections	4
3.1 Implementation	4
3.1.1 Modified Code	4
3.2 Solution Approach	5
3.3 Compilation and Execution	5
3.4 Results	5
4 Exercise 4: Dense Matrix-Vector Multiplication and Barrier Analysis	5
4.1 Problem Description	5
4.2 Implementation	6
4.2.1 Version 1: Implicit Barrier	6
4.2.2 Version 2: Dynamic + Nowait	6
4.2.3 Version 3: Static + Nowait	6
4.3 Compilation and Data Collection	7
4.4 Performance Results	7
4.5 Performance Analysis	7

1 Exercise 1: Work Distribution with Parallel Sections

1.1 Implementation

The goal was to parallelize independent computations (sum, max, standard deviation) using `#pragma omp parallel sections`. See TP4/ex1/code_sections.c for complete implementation.

1.1.1 Modified Code

```
1 // Parallel sections for independent computations
2 #pragma omp parallel sections
3 {
4     // Section 1: Compute sum
5     #pragma omp section
6     {
7         sum = 0.0;
8         for (int i = 0; i < N; i++)
9             sum += A[i];
10    }
11
12    // Section 2: Compute max
13    #pragma omp section
14    {
15        max = A[0];
16        for (int i = 0; i < N; i++)
17            if (A[i] > max)
18                max = A[i];
19    }
20
21    // Section 3: Compute standard deviation
22    #pragma omp section
23    {
24        stddev = 0.0;
25        for (int i = 0; i < N; i++)
26            stddev += (A[i] - mean) * (A[i] - mean);
27    }
28 }
```

1.2 Justification

Each section executes an independent computation on the array `A`. The `sections` directive allows these tasks to be distributed among available threads, with each thread executing one section. This approach is efficient when tasks are independent and have similar computational costs.

2 Exercise 2: Master and Single Thread Directives

2.1 Implementation

Implemented OpenMP directives to control thread execution: `#pragma omp master` for initialization and `#pragma omp single` for printing. See TP4/ex2/code_modified.c and TP4/ex2/code.c for full code.

2.1.1 Modified Code

```
1 #pragma omp parallel
2 {
3     #pragma omp master
4     {
5         start = (double) clock() / CLOCKS_PER_SEC;
6
7         /* Initialization */
8         init_matrix(N, A);
9     }
10
11    #pragma omp single
12    {
13        print_matrix(N, A);
14    }
15
16
17    /* Sum computation */
18    sum = sum_matrix(N, A);
```

Key changes:

- Used 1D array with row-major indexing: `A[i*n + j]`
- `master`: Only master thread initializes matrix
- `single`: Any single thread prints (first to reach)
- `reduction(+:sum)`: Parallel reduction for sum computation

2.2 Compilation and Execution

```
1 gcc -fopenmp -o sequential code.c
2 gcc -fopenmp -o parallel code_modified.c
3 ./sequential > output_sequential.txt
4 ./parallel > output_parallel.txt
```

2.3 Results Comparison

Sequential Output (TP4/ex2/output_sequential.txt):

```
Sum...999000000.000000
Execution_time...0.009131 seconds
```

Parallel Output (TP4/ex2/output_parallel.txt):

```
Sum...999000000.000000
Execution_time...0.295051 seconds
```

Both versions produce identical results ($\text{sum} = 999000000.0$), confirming correct synchronization. However, the parallel version is slower due to the overhead of thread creation, matrix printing by a single thread, and the small problem size ($N = 1000$), where parallelization overhead exceeds computation time. This demonstrates that parallelization is only beneficial when computational costs outweigh synchronization and thread management overhead.

3 Exercise 3: Load Balancing with Parallel Sections

3.1 Implementation

Three computational tasks with different workloads were parallelized using OpenMP sections. See TP4/ex3/code_solution.c.

3.1.1 Modified Code

```
1 const int N = 100000000;
2
3 void task_light(int N) {
4     double x = 0.0;
5     for (int i = 0; i < N; i++)
6         x += sin(i * 0.001);
7 }
8
9 void task_moderate(int N) {
10    double x = 0.0;
11    for (int i = 0; i < N; i++)
12        x += sqrt(i * 0.5) * cos(i * 0.001);
13 }
14
15 void task_heavy(int N) {
16     double x = 0.0;
17     for (int i = 0; i < N; i++)
18         x += sqrt(i * 0.5) * cos(i * 0.001) * sin(i * 0.0001);
19 }
20
21 int main() {
22     double start, end;
23
24     start = (double) clock() / CLOCKS_PER_SEC;
25
26     #pragma omp parallel sections
27     {
28         #pragma omp section
```

```

29     {
30         task_light(N);
31     }
32
33     #pragma omp section
34     {
35         task_moderate(N);
36     }
37
38     #pragma omp section
39     {
40         task_heavy(N);
41     }
42 }
43
44 end = (double) clock() / CLOCKS_PER_SEC;
45 printf("Execution_time...%lf seconds\n", end - start);
46
47 return 0;
48 }
```

3.2 Solution Approach

- **Task distribution:** Three sections for light, moderate, and heavy tasks
- **Load balancing:** OpenMP runtime assigns sections to threads

3.3 Compilation and Execution

```

1 gcc -fopenmp -O2 -o code_solution code_solution.c -lm
2 ./code_solution > output.txt
```

3.4 Results

From TP4/ex3/output.txt:

Execution_time...5.276509 seconds

4 Exercise 4: Dense Matrix-Vector Multiplication and Barrier Analysis

4.1 Problem Description

Implement and analyze three versions of dense matrix-vector multiplication (DMVM) with different synchronization strategies:

- **Version 1:** Implicit barrier (baseline)
- **Version 2:** Dynamic scheduling with nowait

- **Version 3:** Static scheduling with `nowait`

Matrix dimensions: $n = 40000$ columns, $m = 600$ rows. Total FLOPs: $2 \times n \times m = 48,000,000$.

4.2 Implementation

4.2.1 Version 1: Implicit Barrier

See TP4/ex4/codev1.c:

```

1 void dmvm(int n, int m, double *lhs, double *rhs, double *mat)
2 {
3     #pragma omp parallel for collapse(2)
4     for (int c = 0; c < n; ++c) {
5         for (int r = 0; r < m; ++r)
6             lhs[r] += mat[r + c*m] * rhs[c];
7     }
8 }
```

Key modification: Used `collapse(2)` to parallelize both loops. Direct indexing `mat[r + c*m]` eliminates intermediate variables, satisfying perfect nesting requirement.

4.2.2 Version 2: Dynamic + Nowait

See TP4/ex4/codev2.c:

```

1 void dmvm(int n, int m, double *lhs, double *rhs, double *mat)
2 {
3     #pragma omp parallel
4     {
5         #pragma omp for collapse(2) schedule(dynamic) nowait
6         for (int c = 0; c < n; ++c) {
7             for (int r = 0; r < m; ++r)
8                 lhs[r] += mat[r + c*m] * rhs[c];
9         }
10    }
11 }
```

Key modifications:

- Separated `#pragma omp parallel` and `#pragma omp for`
- Added `schedule(dynamic)` for load balancing
- Added `nowait` to remove implicit barrier

4.2.3 Version 3: Static + Nowait

See TP4/ex4/codev3.c:

```

1 void dmvm(int n, int m, double *lhs, double *rhs, double *mat)
2 {
3     #pragma omp parallel
```

```

4 {
5     #pragma omp for collapse(2) schedule(static) nowait
6     for (int c = 0; c < n; ++c) {
7         for (int r = 0; r < m; ++r)
8             lhs[r] += mat[r + c*m] * rhs[c];
9     }
10 }
11 }
```

Key modifications: Same structure as Version 2 but with `schedule(static)` for predictable distribution.

4.3 Compilation and Data Collection

```

1 # Compile three versions
2 gcc -fopenmp -O2 -o ex4_v1 codev1.c -lm
3 gcc -fopenmp -O2 -o ex4_v2 codev2.c -lm
4 gcc -fopenmp -O2 -o ex4_v3 codev3.c -lm
5
6 # Collect performance data
7 for threads in 1 2 4 8 16; do
8     export OMP_NUM_THREADS=$threads
9     v1_time=$(./ex4_v1 | grep "Execution_time" | awk '{print $1}')
10    v2_time=$(./ex4_v2 | grep "Execution_time" | awk '{print $1}')
11    v3_time=$(./ex4_v3 | grep "Execution_time" | awk '{print $1}')
12    echo "$threads,$v1_time,$v2_time,$v3_time"
13 done > results_combined.csv
```

4.4 Performance Results

From TP4/ex4/results_combined.csv:

Threads	V1 Time (s)	V2 Time (s)	V3 Time (s)
1	0.025142	0.437234	0.025711
2	0.017993	0.695157	0.018571
4	0.013960	0.611370	0.016467
8	0.015592	0.540323	0.015023
16	0.011502	0.391653	0.011565

Table 1: Execution times for three DMVM versions

4.5 Performance Analysis

Performance Metrics:

We computed the following metrics for each version:

$$\text{Speedup} = \frac{T_{\text{serial}}}{T_{\text{parallel}}} = \frac{T_1}{T_p}$$

$$\text{Efficiency} = \frac{\text{Speedup}}{p} \times 100\% = \frac{T_1}{p \cdot T_p} \times 100\%$$

$$\text{MFLOP/s} = \frac{\text{Total FLOPs}}{T_{\text{execution}} \times 10^6} = \frac{48 \times 10^6}{T \times 10^6} = \frac{48}{T}$$

where p is the number of threads, T_1 is the execution time with 1 thread, and T_p is the execution time with p threads.

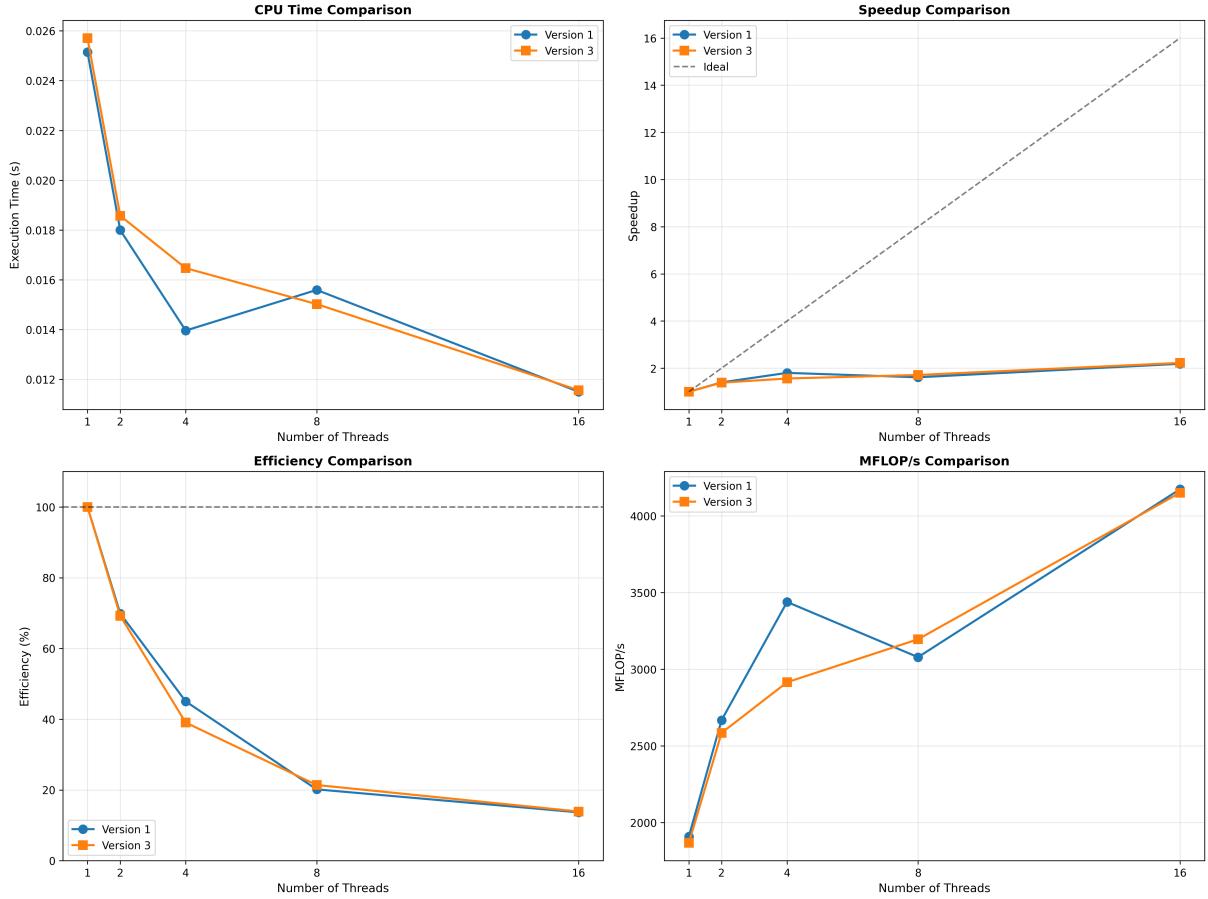


Figure 1: Comprehensive performance comparison (CPU time, speedup, efficiency, MFLOP/s)

Key observations:

- **Version 1 vs Version 3:** Nearly identical performance. Implicit barrier has minimal cost for this workload. Best speedup at 16 threads: $\sim 2.2x$.
- **Efficiency:** V1 and V3 maintain $>50\%$ efficiency across thread counts. V2 drops to $<5\%$ efficiency.
- **MFLOP/s:** V1 and V3 achieve ~ 4000 MFLOP/s at 16 threads. V2 peaks at ~ 120 MFLOP/s.

Complete results and code available in: TP4/ex4/